


1

Relational Databases

Introduction

- Relational database management systems (RDBMS) were proposed in the 1970s by E.F. Codd
- This paradigm dominated database design until the turn of the century
- Examples of RDBMS include:
 - **MySQL**
 - MS SQL
 - MS Access
 - Oracle SQL



2

Relational Databases

Introduction

- Relational database management systems (RDBMS) have a number of defining characteristics
 - Data in a RDBMS is organised through the “Relational Model”
 - Support Structured Query Language (SQL) for interacting with a database
 - Support relationships between tables
 - Typically support ACID transactions



3

Relational Databases

Relational Model


- Relational models are essentially tables within a database
- Tables generally represent one entity type, such as student/customer records
- These tables consist of rows and columns
 - Rows may be referred to as Tuples or Records
 - Columns may be referred to as Fields or Attributes
 - Tables may be referred to as Relations or Base Revlar*



4

Relational Databases

Relational Model



5

Relational Databases

Relational Model

| Customer ID | Tax ID | Name | Address | [More fields...] |
|-------------|-------------|------------|------------------|------------------|
| 1234567890 | 555-5512222 | Munmun | 323 Broadway | ... |
| 2223344556 | 555-5523232 | Wile E. | 1200 Main Street | ... |
| 3334445563 | 555-5533323 | Ekta | 871 1st Street | ... |
| 4232342432 | 555-5325523 | E. F. Codd | 123 It Way | ... |



6

Relational Databases

Relational Model

- The relational model can have a number of keys to identify records, these include:
 - Primary Keys (PK)
 - Composite Keys (CK)
 - Foreign Keys (FK)



7

Relational Databases

Relational Model

- Primary Keys
 - A field offering a unique identifier
 - Used to identify a unique tuple
 - May be an auto-incrementing numeric value

| Store ID | Purchase Location |
|----------|-------------------|
| 1 | Los Angeles |
| 2 | New York |
| 3 | San Francisco |



8

Relational Databases

Relational Model

- Composite Keys
 - Used to identify a unique tuple
 - A combination of two or more columns
 - May employ auto incrementing on one, some or all columns involved

| Customer ID | Store ID | Purchase Location |
|-------------|----------|-------------------|
| 1 | 1 | Los Angeles |
| 1 | 3 | San Francisco |
| 2 | 1 | Los Angeles |
| 3 | 2 | New York |
| 4 | 3 | San Francisco |



9

Relational Databases

Relational Model

- Foreign Keys
 - A field in one table that uniquely identifies a key in another table
 - Facilitates selection of records across tables
 - Enforces referential integrity
 - May be an auto-incrementing numeric value



10

Relational Databases

Relational Model

- Foreign Keys
 - A field in one table that uniquely identifies a key in another table
 - Facilitates selection of records across tables
 - Enforces referential integrity
 - A FK may only contain values from a specified column in a foreign table or a null value – providing assurance that related information exists
 - When deleting a record containing a FK the RDBMS may delete the record in the other table or prevent deletion

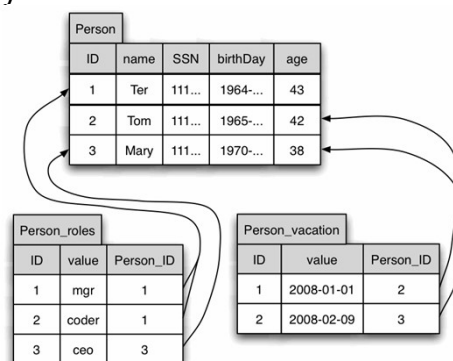


11

Relational Databases

Relational Model

- Foreign Keys



12

Relational Databases

Normalization

- Relations/Tables should be carefully defined
- The databases should be “Normalized”
- Normalization is the process of reorganising attributes into tables with the goal of minimising redundancy
 - Increases storage efficiency
 - May have performance implications
- Normalization may take a number of forms



13

Relational Databases

Normalization

- An example of normalization in a sale database
 - Storing sales records and product records in separate tables
 - Product records are stored within unique IDs
 - Records in the sales table reference items sold through referencing a unique ID



14

Relational Databases

Normalization

- Three different levels of normalization
 - 1NF – First Normal Form
 - 2NF – Second Normal Form
 - 3NF – Third Normal Form
- A DB is described as normalized when it satisfies the 3NF



15

Relational Databases

Normalization – 1NF

- 1NF enforces the following criteria
 - Eliminate repetition in individual tables
 - Create a separate table for each set of related data/entity type
 - Identify each record with a primary key
 - Contains only atomic values



16

Relational Databases

Normalization – 1NF

- Given the requirement to store multiple phone numbers
- A violates the requirement as telephone number may only contain one atomic value
- B satisfies the 1NF but not more complex forms

| Customer ID | First Name | Surname | Telephone Number |
|-------------|------------|-----------|------------------------------|
| 123 | Robert | Ingram | 555-861-2025 |
| 456 | Jane | Wright | 555-403-1659 555-776-4100 |
| 789 | Maria | Fernandez | 555-808-9633 |

A

| Customer ID | First Name | Surname | Telephone Number |
|-------------|------------|-----------|------------------|
| 123 | Robert | Ingram | 555-861-2025 |
| 456 | Jane | Wright | 555-403-1659 |
| 456 | Jane | Wright | 555-776-4100 |
| 789 | Maria | Fernandez | 555-808-9633 |

B



17

Relational Databases

Normalization – 2NF

- 2NF enforces the following criteria
 - Satisfy criteria expressed in 1NF
 - Non-key attributes are fully functionally dependent on the primary key



18

Relational Databases

Normalization – 2NF

- A presents a table using a CK that is not normalised.
- A violates the requirement as Purchase Location only depends on store ID which is only an element of the PK
- B satisfies the 2NF

| Customer ID | Store ID | Purchase Location |
|-------------|----------|-------------------|
| 1 | 1 | Los Angeles |
| 1 | 3 | San Francisco |
| 2 | 1 | Los Angeles |
| 3 | 2 | New York |
| 4 | 3 | San Francisco |

A

| Customer ID | Store ID |
|-------------|----------|
| 1 | 1 |
| 1 | 3 |
| 2 | 1 |
| 3 | 2 |
| 4 | 3 |

B



19

Relational Databases

Normalization – 3NF

- 3NF enforces the following criteria
 - Satisfy criteria expressed in 2NF
 - “Every non-prime attribute of R is non-transitively dependent on every key of R”
- More simply put:
“Every non-key attribute must provide a fact about the key, the whole key, and nothing but the key.” – Bill Kent



20

Relational Databases

Normalization – 3NF

- A is not in the 3NF, Book ID determines Genre ID and Genre ID determines Genre Type. This represents a transitive functional dependency.

| Book ID | Genre ID | Genre Type | Price |
|---------|----------|------------|-------|
| 1 | 1 | Gardening | 25.99 |
| 2 | 2 | Sports | 14.99 |
| 3 | 1 | Gardening | 10.00 |
| 4 | 3 | Travel | 12.99 |
| 5 | 2 | Sports | 17.99 |

A

- B is brought to the 3rd normal form by splitting into two tables

| Book ID | Genre ID | Price |
|---------|----------|-------|
| 1 | 1 | 25.99 |
| 2 | 2 | 14.99 |
| 3 | 1 | 10.00 |
| 4 | 3 | 12.99 |
| 5 | 2 | 17.99 |

| Genre ID | Genre Type |
|----------|------------|
| 1 | Gardening |
| 2 | Sports |
| 3 | Travel |

B



21

Relational Databases

Entity-Relationship Modelling

- Mapping between tables may take a number of forms:
 - 1 to 1 – e.g one person may have one social security number
 - 1 to Many - e.g one person may have many bank accounts
 - Many to Many – e.g. books may have many authors, and vice versa

| Book ID | Genre ID | Price |
|---------|----------|-------|
| 1 | 1 | 25.99 |
| 2 | 2 | 14.99 |
| 3 | 1 | 10.00 |
| 4 | 3 | 12.99 |
| 5 | 2 | 17.99 |

| Genre ID | Genre Type |
|----------|------------|
| 1 | Gardening |
| 2 | Sports |
| 3 | Travel |



22

Relational Databases

ACID transactions

- RDBMS typically support ACID transactions
- ACID transactions offer the following properties
 - Atomicity – an entire operation is performed or nothing is
 - Consistency – any transaction will bring the DB from one valid state to another
 - Isolation – concurrent transactions would not interfere with one another
 - Durability – once a transaction has been committed it will remain even in the event of power loss or and error



23

Relational Databases

SQL

- SQL is the primary mechanism for interacting with RDBMS
- SQL is subdivided into a number of language elements, including:
 - Clauses - components of statements and queries
 - Expressions – used to evaluate a value
 - Predicates – conditions to be evaluated within SQL
 - Queries – Retrieve data based on criteria
 - Statements – instructions which may have an effect on data or DB schema



24

Relational Databases

SQL

UPDATE clause { UPDATE country
 SET clause { SET population = population + 1
 WHERE clause { WHERE name = 'USA';

Expression
 Expression
 Predicate
 Statement



25

Relational Databases

SQL

- Examples of clauses include:
 - SELECT – select data from a database
 - INSERT – insert data into a database
 - UPDATE – update records in a database
 - DELETE – delete records in a database
 - WHERE – used to restrict the operation of a statement
 - LIMIT – used to limit the number of records returned by a resultset



26

Relational Databases

SQL

- Examples of clauses include
 - JOIN – used to combine rows from 2+ tables
 - DISTINCT – used to remove duplicates from the result set produced by a select statement
 - ORDER BY – used to order the result set of a query by a specified column list
 - GROUP BY – used in a select statement to collect data across multiple records and group the results



27

Relational Databases

SQL

- Examples of clauses include
 - HAVING – similar to where but applied to aggregate functions, such as GROUP BY
 - UNION - combine results of 2+ select statements without returning duplicate rows



28

Relational Databases

SQL

- The following is the structure for a statement that creates a table within SQL

```
CREATE TABLE table_name
(
  column_name1 data_type(size),
  column_name2 data_type(size),
  column_name3 data_type(size),
  PRIMARY KEY (column_name)
);
```



29

Relational Databases

SQL – Data types

- SQL supports a number of data types including:
 - CHAR – a number of characters up to 255; char (30)
 - VARCHAR – a number of characters up to 65,535; varchar (300)
 - TEXT – strings/text cannot be ordered or used in a where clause, in most cases
 - BINARY – binary values
 - BLOB – binary strings



30

Relational Databases

SQL – Data types

- SQL supports a number of data types including:
 - INT – whole numbers, 4 bytes, signed or unsigned
 - -2147483648 to 2147483647
 - 0 to 4294967295
 - BIGINT – whole numbers, 8 bytes, signed or unsigned
 - -9223372036854775808 to 9223372036854775807
 - 0 to 18446744073709551615
- These may vary with RDBMS implementation
- Correct selection and specification is imperative
- Incorrect specification has performance implications



31

Relational Databases

Storage hierarchy

There is a storage hierarchy present within Relational databases:

1. Database server
2. Schemas/Databases*
3. Tables
4. Relations



32

Relational Databases

Schema

Schemas/databases are used to store related tables.

The SQL command to create a schema is:

```
CREATE SCHEMA `example_schema` ;
```

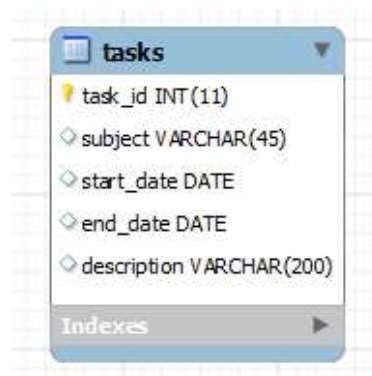
Once a schema is in place, tables may be created.



33

Relational Databases

Tables



34

Relational Databases

Tables

- To create this table the following statement is used:

```
CREATE TABLE IF NOT EXISTS example_schema.tasks (
  task_id INT(11) NOT NULL AUTO_INCREMENT,
  subject VARCHAR(45) DEFAULT NULL,
  start_date DATE DEFAULT NULL,
  end_date DATE DEFAULT NULL,
  description VARCHAR(200) DEFAULT NULL,
  PRIMARY KEY (task_id)
) ENGINE=InnoDB
```



35

Relational Databases

Relations/Records

- After a table is created a relation/record may inserted
- To insert a relation/record the following SQL is used

```
INSERT INTO `example_schema`.`tasks`
(`subject`, `start_date`, `end_date`, `description`)
VALUES
('SubjectName', '1970-01-01', '1970-01-01', 'A lengthy
description');
```

Note the missing field



36

Relational Databases

Relations/Records

- After a relation is inserted it may be Selected as part of a result set

```
SELECT * FROM example_schema.tasks
WHERE task_id > 0 ORDER BY subject;
```

| task_id | subject | start_date | end_date | description |
|---------|--------------|------------|------------|-----------------|
| 2 | ASubjectName | 1970-01-01 | 1970-02-01 | Another desc |
| 1 | SubjectName | 1970-01-01 | 1970-01-01 | A lengthy de... |
| ▶* | NULL | NULL | NULL | |



37

Relational Databases

Relations/Records

- Records present within an database may be updated through the following command

```
UPDATE example_schema.tasks
SET subject="UpdatedSubject" WHERE task_id = 2;
```

| task_id | subject | start_date | end_date | description |
|---------|----------------|------------|------------|-----------------|
| 1 | SubjectName | 1970-01-01 | 1970-01-01 | A lengthy de... |
| 2 | UpdatedSubject | 1970-01-01 | 1970-02-01 | Another desc |
| ▶* | NULL | NULL | NULL | NULL |



38

Relational Databases

Relations/Records

- Records may be deleted with the DELETE command

```
DELETE FROM example_schema.tasks
WHERE task_id > 1;
```

| | | | | |
|---|-------------|------------|------------|-----------------|
| 1 | SubjectN... | 1970-01-01 | 1970-01-01 | A lengthy de... |
| * | NULL | NULL | NULL | NULL |



39

Relational Databases

Advanced topics

Join – This allows querying across tables to produce a record set/Revlar

Consider the tasks table coexisting with the following “task_operator” table

| Column | Type | Default Value | Nullable |
|---------------|--------------|---------------|----------|
| operator_id | int(11) | | NO |
| task_id | int(11) | | YES |
| operator_name | varchar(450) | | YES |



40

Relational Databases

Advanced topics

The task operator table contains the following

| | operator_id | task_id | operator_name |
|----|-------------|---------|---------------|
| | 1 | 1 | Anon |
| | 2 | 2 | Bnon |
| | 3 | 3 | Cnon |
| | 4 | 1 | Dnon |
| ▶* | | NULL | NULL |



41

Relational Databases

Advanced topics

Using join it is possible to find the operators associated with a task simply by specifying the task name.

```
SELECT T1.operator_name FROM example_schema.task_operator T1
INNER JOIN example_schema.tasks T2 ON T2.task_id = T1.task_id
WHERE T2.subject = "SubjectName";
```

| | operator_name |
|--|---------------|
| | Anon |
| | Dnon |



42

Relational Databases

Advanced topics

DISTINCT – when distinct is used in a clause it will select remove duplicates.

E.G `SELECT DISTINCT(SubjectName) FROM example_schema.tasks;`

Stored Procedures – SQL statements stored within a DB, removes a layer of logic from consuming applications.

Triggers – A procedure or function stored within a DB that activates the a certain condition is met.



43

Relational Databases

Use cases

- Use for relational records where
 - Single / low numbers of related instances will be stored*
 - Applications where ACID consistency is required
- Examples include
 - User profiles, e.g. name + login info
 - Reliable storage of transactional records, such as financial logs**
 - Address books



44

Relational Databases

Use cases

- When not to use a relational database
 - High volume data, such as logs that do not require ACID consistency
 - Storage of binary data; such as images, videos and objects*
 - Storage of XML**



45

Relational Databases

- Break time
- Practical aspect to follow



46

